

Microsoft Malware Data Challenge

As one part of their overall strategy for doing so, Microsoft is challenging the data science community to develop techniques to predict if a machine will soon be hit with malware.

<https://www.kaggle.com/c/microsoft-malware-prediction> (<https://www.kaggle.com/c/microsoft-malware-prediction>)

```
In [168]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.pyplot import figure
import seaborn as sns
```

1 Data Setup and Exploration

1.1 Loading MSFT Data

```

In [169]: # referred https://www.kaggle.com/theoviel/load-the-totality-of-the-data
dtypes = {
    'MachineIdentifier': 'category',
    'ProductName': 'category',
    'EngineVersion': 'category',
    'AppVersion': 'category',
    'AvSigVersion': 'category',
    'IsBeta': 'int8',
    'RtpStateBitfield': 'float16',
    'IsSxsPassiveMode': 'int8',
    'DefaultBrowsersIdentifier': 'float32',
    'AVProductStatesIdentifier': 'float32',
    'AVProductsInstalled': 'float16',
    'AVProductsEnabled': 'float16',
    'HasTpm': 'int8',
    'CountryIdentifier': 'int16',
    'CityIdentifier': 'float32',
    'OrganizationIdentifier': 'float16',
    'GeoNameIdentifier': 'float16',
    'LocaleEnglishNameIdentifier': 'int16',
    'Platform': 'category',
    'Processor': 'category',
    'OsVer': 'category',
    'OsBuild': 'int16',
    'OsSuite': 'int16',
    'OsPlatformSubRelease': 'category',
    'OsBuildLab': 'category',
    'SkuEdition': 'category',
    'IsProtected': 'float16',
    'AutoSampleOptIn': 'int8',
    'PuaMode': 'category',
    'SMode': 'float16',
    'IeVerIdentifier': 'float16',
    'SmartScreen': 'category',
    'Firewall': 'float16',
    'UacLuaenable': 'float32',
    'UacLuaenable': 'float64',
    'Census_MDC2FormFactor': 'category',
    'Census_DeviceFamily': 'category',
    'Census_OEMNameIdentifier': 'float32',
    'Census_OEMModelIdentifier': 'float32',
    'Census_ProcessorCoreCount': 'float16',
    'Census_ProcessorManufacturerIdentifier': 'float16',
    'Census_ProcessorModelIdentifier': 'float32',
    'Census_ProcessorClass': 'category',
    'Census_PrimaryDiskTotalCapacity': 'float64',
    'Census_PrimaryDiskTypeName': 'category',
    'Census_SystemVolumeTotalCapacity': 'float64',
    'Census_HasOpticalDiskDrive': 'int8',
    'Census_TotalPhysicalRAM': 'float32',
    'Census_ChassisTypeName': 'category',
    'Census_InternalPrimaryDiagonalDisplaySizeInInches': 'float32',
    'Census_InternalPrimaryDisplayResolutionHorizontal': 'float32',
    'Census_InternalPrimaryDisplayResolutionVertical': 'float32',
    'Census_PowerPlatformRoleName': 'category',
    'Census_InternalBatteryType': 'category',

```

```

'Census_InternalBatteryNumberOfCharges': 'float64',
'Census_OSVersion': 'category',
'Census_OSArchitecture': 'category',
'Census_OSBranch': 'category',
'Census_OSBuildNumber': 'int16',
'Census_OSBuildRevision': 'int32',
'Census_OSEdition': 'category',
'Census_OSSkuName': 'category',
'Census_OSInstallTypeName': 'category',
'Census_OSInstallLanguageIdentifier': 'float16',
'Census_OSUILocaleIdentifier': 'int16',
'Census_OSWUAutoUpdateOptionsName': 'category',
'Census_IsPortableOperatingSystem': 'int8',
'Census_GenuineStateName': 'category',
'Census_ActivationChannel': 'category',
'Census_IsFlightingInternal': 'float16',
'Census_IsFlightsDisabled': 'float16',
'Census_FlightRing': 'category',
'Census_ThresholdOptIn': 'float16',
'Census_FirmwareManufacturerIdentifier': 'float16',
'Census_FirmwareVersionIdentifier': 'float32',
'Census_IsSecureBootEnabled': 'int8',
'Census_IsWIMBootEnabled': 'float16',
'Census_IsVirtualDevice': 'float16',
'Census_IsTouchEnabled': 'int8',
'Census_IsPenCapable': 'int8',
'Census_IsAlwaysOnAlwaysConnectedCapable': 'float16',
'Wdft_IsGamer': 'float16',
'Wdft_RegionIdentifier': 'float16',
'HasDetections': 'int8'
}

```

```
In [279]: train = pd.read_csv('MSFT_DATA/train_sample_50k.csv', dtype=dtypes)
```

1.2 Data Exploration

```
In [280]: print("%d Columns" % (len(train.columns)))  
print("%d Rows" % (len(train)))  
train.head()
```

83 Columns
50000 Rows

Out[280]:

	MachineIdentifier	ProductName	EngineVersion	AppVersion	AvSigVersion
0	0000028988387b115f69f31a3bf04f09	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1735.0
1	000007535c3f730efa9ea0b7ef1bd645	win8defender	1.1.14600.4	4.13.17134.1	1.263.48.0
2	000007905a28d863f6d0d597892cd692	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1341.0
3	00000b11598a75ea8ba1beea8459149f	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1527.0
4	000014a5f00daa18e76b81417eeb99fc	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1379.0

5 rows × 83 columns

```
In [281]: sum([True in i for i in train.isnull().values])
```

Out[281]: 50000

```
In [282]: sum(['nan' in i for i in train.values])
```

Out[282]: 0


```
In [283]: figure(num=None, figsize=(8, 6), dpi=80, facecolor='w', edgecolor='k')

plt.bar(["Categorical Regular", "Numerical Continuous", "Numerical Discrete"])

plt.title('Counts of Feature Variable Types')
plt.xlabel('Variable Type')
plt.ylabel('Count')

plt.savefig('Var-Type-Bar.png')

plt.show()
```



2 Feature Engineering

2.1 Missing Entries

```
In [284]: bad_features = []
```

```
In [285]: # List of percentage of null values...
(train.isnull().sum()/train.shape[0]).sort_values(ascending=False)
```

Census_IsTouchEnabled	0.0000
Census_IsSecureBootEnabled	0.0000
Census_FlightRing	0.0000
Census_ActivationChannel	0.0000
Census_GenuineStateName	0.0000
Census_IsPortableOperatingSystem	0.0000
Census_OSWUAutoUpdateOptionsName	0.0000
Census_OSUILocaleIdentifier	0.0000
Census_OSInstallTypeName	0.0000
Census_OSSkuName	0.0000
Census_OSEdition	0.0000
Census_OSBuildRevision	0.0000
Census_OSBuildNumber	0.0000
Census_OSBranch	0.0000
Census_OSArchitecture	0.0000
Census_OSVersion	0.0000
Census_PowerPlatformRoleName	0.0000
Census_HasOpticalDiskDrive	0.0000
Census_DeviceFamily	0.0000
Census_MDC2FormFactor	0.0000

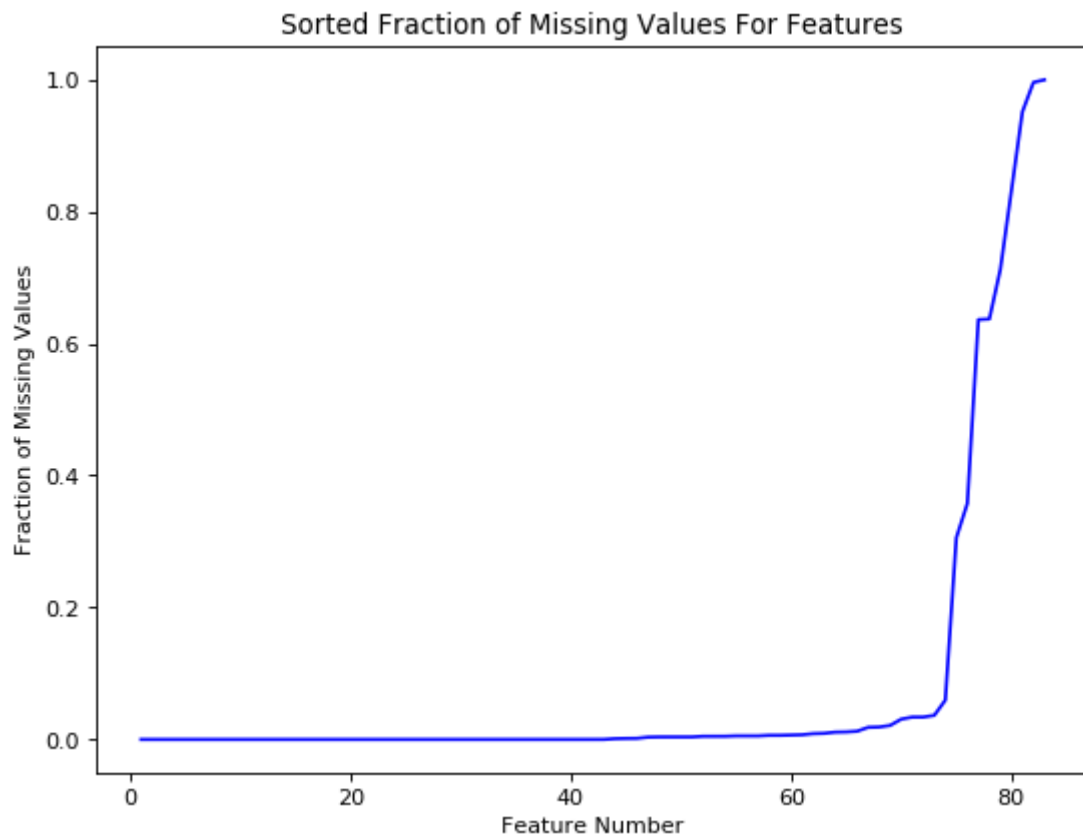
```
In [286]: missing_vals = (train.isnull().sum()/train.shape[0]).sort_values(ascending=False)
```

```
In [287]: figure(num=None, figsize=(8, 6), dpi=80, facecolor='w', edgecolor='k')

plt.plot(list(range(1, len(missing_vals) + 1, 1)), missing_vals, color='blue')

# Plot text
plt.title('Sorted Fraction of Missing Values For Features')
plt.xlabel('Feature Number')
plt.ylabel('Fraction of Missing Values')

plt.savefig('Missing_Features.png', dpi=420)
plt.show()
```



```
In [288]: bad_features.append('PuaMode')
          bad_features.append('Census_ProcessorClass')
```

2.2 Skewed Features

For categorical data, extremely skewed data means most of the datapoints belong to a particular category for that feature.

i.e. House color but all the houses in the data are red. (Just remove house color)

```
In [289]: pd.options.display.float_format = '{:,.4f}'.format

sk_df = pd.DataFrame([{'column': c, 'uniq': train[c].nunique(),
                      'skewness': train[c].value_counts(normalize=True)
                      .values[0]} for c in train.columns])

# Skewness
sk_df = sk_df.sort_values('skewness', ascending=False)
sk_df.head()
```

Out[289]:

	column	skewness	uniq
27	AutoSampleOptIn	1.0000	1
28	PuaMode	1.0000	1
69	Census_IsFlightsDisabled	1.0000	1
5	IsBeta	1.0000	1
75	Census_IsWIMBootEnabled	1.0000	1

In [290]: sk_df

Out[290]:

	column	skewness	uniq
27	AutoSampleOptIn	1.0000	1
28	PuaMode	1.0000	1
69	Census_IsFlightsDisabled	1.0000	1
5	IsBeta	1.0000	1
75	Census_IsWIMBootEnabled	1.0000	1
68	Census_IsFlightingInternal	1.0000	1
71	Census_ThresholdOptIn	0.9998	2
29	SMode	0.9996	2
65	Census_IsPortableOperatingSystem	0.9992	2
35	Census_DeviceFamily	0.9987	2
33	UacLuaenable	0.9936	3
76	Census_IsVirtualDevice	0.9932	2
1	ProductName	0.9888	2
12	HasTpm	0.9874	2
7	IsSxsPassiveMode	0.9823	2
32	Firewall	0.9772	2
11	AVProductsEnabled	0.9748	5
6	RtpStateBitfield	0.9734	6
20	OsVer	0.9679	8
18	Platform	0.9666	4
78	Census_IsPenCapable	0.9624	2
26	IsProtected	0.9464	2
79	Census_IsAlwaysOnAlwaysConnectedCapable	0.9420	2
70	Census_FlightRing	0.9360	7
45	Census_HasOpticalDiskDrive	0.9241	2
19	Processor	0.9098	3
55	Census_OSArchitecture	0.9098	3
66	Census_GenuineStateName	0.8837	4
39	Census_ProcessorManufacturerIdentifier	0.8816	3
77	Census_IsTouchEnabled	0.8737	2
...
57	Census_OSBuildNumber	0.4503	31
64	Census_OSWUAutoUpdateOptionsName	0.4407	6

	column	skewness	uniq
21	OsBuild	0.4392	33
23	OsPlatformSubRelease	0.4392	9
30	leVerIdentifier	0.4385	132
2	EngineVersion	0.4331	38
24	OsBuildLab	0.4110	316
59	Census_OSEdition	0.3885	18
60	Census_OSSkuName	0.3885	16
62	Census_OSInstallLanguageIdentifier	0.3586	39
63	Census_OSUILocaleIdentifier	0.3554	55
48	Census_InternalPrimaryDiagonalDisplaySizeInches	0.3416	304
42	Census_PrimaryDiskTotalCapacity	0.3171	318
72	Census_FirmwareManufacturerIdentifier	0.3123	136
61	Census_OSInstallTypeName	0.2945	9
17	LocaleEnglishNameIdentifier	0.2330	171
81	Wdft_RegionIdentifier	0.2086	15
16	GeoNameIdentifier	0.1720	226
58	Census_OSBuildRevision	0.1561	209
54	Census_OSVersion	0.1561	241
36	Census_OEMNameIdentifier	0.1448	640
8	DefaultBrowsersIdentifier	0.1096	203
13	CountryIdentifier	0.0445	216
37	Census_OEMModelIdentifier	0.0344	10800
40	Census_ProcessorModelIdentifier	0.0324	1509
4	AvSigVersion	0.0116	3105
14	CityIdentifier	0.0115	10423
73	Census_FirmwareVersionIdentifier	0.0107	8930
44	Census_SystemVolumeTotalCapacity	0.0056	26255
0	MachineIdentifier	0.0000	50000

83 rows × 3 columns

```
In [291]: skew_thresh = .99

bad_features.extend(sk_df[sk_df.skewness > skew_thresh].column.tolist())
bad_features = list(set(bad_features))
bad_features
```

```
Out[291]: ['Census_IsVirtualDevice',
'Census_ProcessorClass',
'AutoSampleOptIn',
'Census_IsFlightingInternal',
'UacLuaenable',
'Census_IsPortableOperatingSystem',
'Census_IsFlightsDisabled',
'Census_IsWIMBootEnabled',
'Census_ThresholdOptIn',
'SMode',
'Census_DeviceFamily',
'PuaMode',
'IsBeta']
```

```
In [292]: # Updating our train data
train.drop(bad_features, axis=1, inplace=True)
```

```
In [293]: # Machine ID shouldn't help and only adds noise since theres too many
train.drop('MachineIdentifier', axis=1, inplace=True)
```

2.3 Purify Missing Data

```
In [294]: sum([True in i for i in train.isnull().values])
```

```
Out[294]: 49540
```

```
In [295]: # Fill in the numerical data
train = train.fillna(train.median())
```

```
In [296]: train.head()
```

```
Out[296]:
```

	ProductName	EngineVersion	AppVersion	AvSigVersion	RtpStateBitfield	IsSxsPassiveMode
0	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1735.0	7.0000	0
1	win8defender	1.1.14600.4	4.13.17134.1	1.263.48.0	7.0000	0
2	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1341.0	7.0000	0
3	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1527.0	7.0000	0
4	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1379.0	7.0000	0

5 rows × 69 columns

```
In [297]: # Fill in the categorical data
train = train.fillna(train.mode().iloc[0])
```

```
In [298]: train.head()
```

```
Out[298]:
```

	ProductName	EngineVersion	AppVersion	AvSigVersion	RtpStateBitfield	IsSxsPassiveMode
0	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1735.0	7.0000	0
1	win8defender	1.1.14600.4	4.13.17134.1	1.263.48.0	7.0000	0
2	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1341.0	7.0000	0
3	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1527.0	7.0000	0
4	win8defender	1.1.15100.1	4.18.1807.18075	1.273.1379.0	7.0000	0

5 rows × 69 columns

2.4 Correlated Features

```
In [299]: from sklearn.preprocessing import LabelEncoder

train['SmartScreen'] = train.SmartScreen.astype('category')
train['Census_InternalBatteryType'] = train.Census_InternalBatteryType.astype('category')

cate_cols = train.select_dtypes(include='category').columns.tolist()

le = LabelEncoder()

for col in cate_cols:
    train[col] = le.fit_transform(train[col])
```

```
from sklearn.preprocessing import LabelEncoder
category_cols = [(train[c].dtype == 'O' or
train[c].dtype == 'int64') for c in train.columns]
le = LabelEncoder()
```

```
train['IsBeta'] = le.fit_transform(train['IsBeta'].astype('category'))
```

```
for col in train.columns[category_cols]:
    try: train[col] = le.fit_transform(train[col].astype('category'))
    except: print(col)
```

```
In [300]: train[train.columns[0:10]].corr()
```

```
Out[300]:
```

	ProductName	EngineVersion	AppVersion	AvSigVersion	RtpStateBitfield	IsSxsPassiveMode
ProductName	1.0000	-0.0128	0.1423	-0.0120	-0.0045	0.0143
EngineVersion	-0.0128	1.0000	0.2581	0.8894	0.0045	-0.0014
AppVersion	0.1423	0.2581	1.0000	0.2398	0.0319	0.1531
AvSigVersion	-0.0120	0.8894	0.2398	1.0000	0.0013	0.0317
RtpStateBitfield	-0.0045	0.0045	0.0319	0.0013	1.0000	-0.0386
IsSxsPassiveMode	0.0143	0.0005	-0.0208	0.0074	-0.8897	
DefaultBrowsersIdentifier	-0.0014	-0.0680	-0.0535	-0.0310	0.0018	
AVProductStatesIdentifier	0.1531	0.0658	0.1108	0.0503	0.1175	
AVProductsInstalled	0.0317	-0.1161	-0.1012	-0.1055	-0.1887	
AVProductsEnabled	-0.0386	-0.0273	-0.0596	-0.0236	-0.0041	

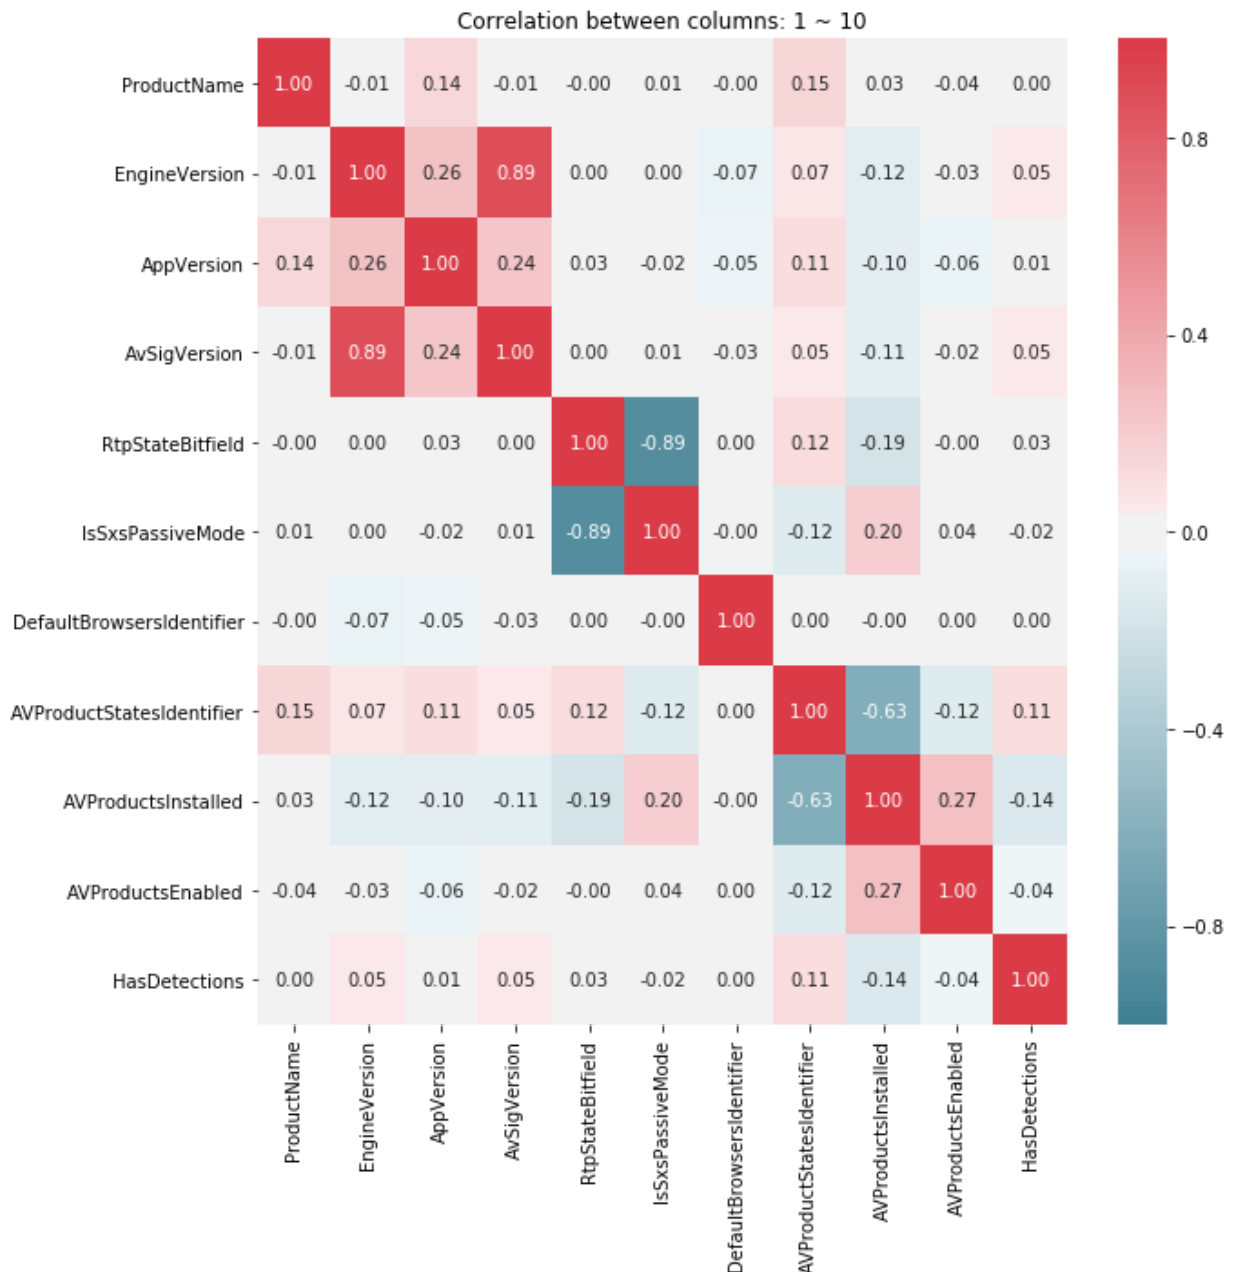
```

In [301]: # Display correlation heat maps

cols = train.columns
for i in range(0, len(cols), 10):
    co_cols = cols[i:i+10].tolist() + ['HasDetections']
    #co_cols.append(pd.Index() 'HasDetections')
    plt.figure(figsize=(10,10))
    colormap = sns.diverging_palette(220, 10, as_cmap=True)
    #Generate Heat Map, allow annotations and place floats in map
    sns.heatmap(train[co_cols].corr(), cmap=colormap, annot=True, fmt=".2f")

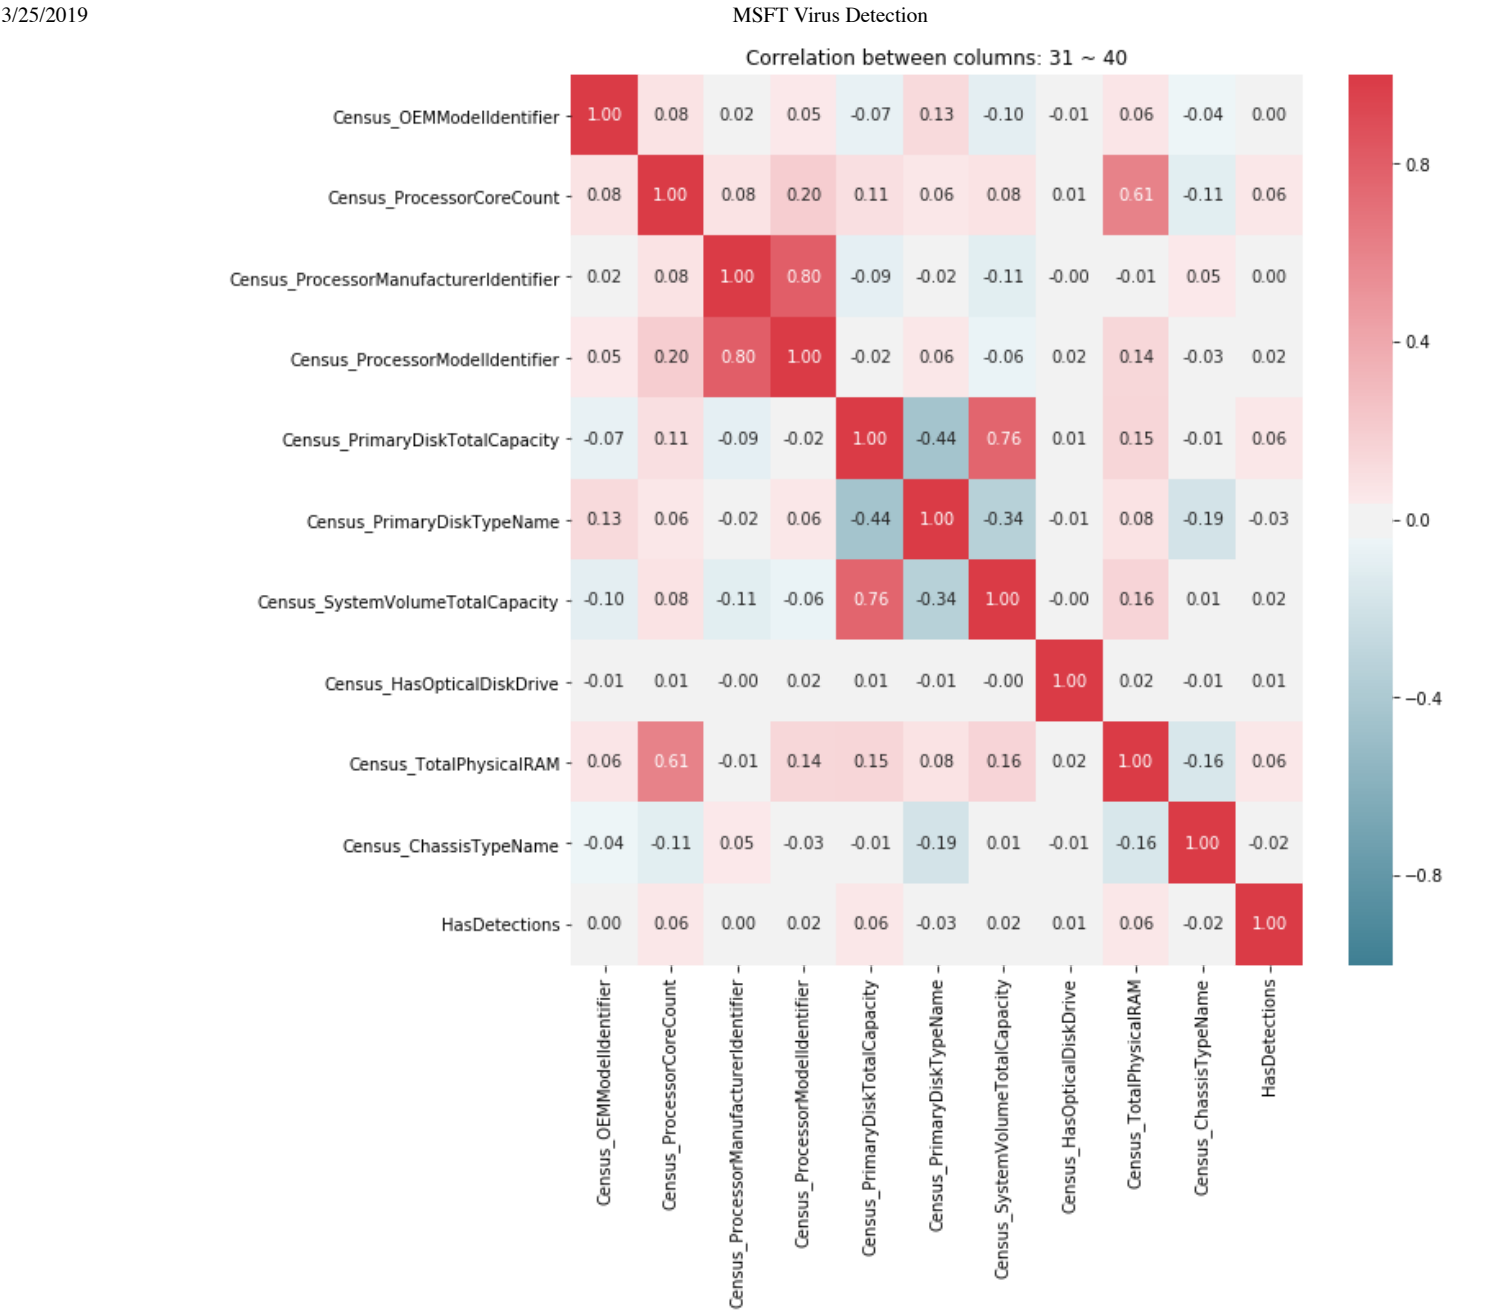
    plt.title('Correlation between columns: %d ~ %d' % (i+1, i+10))
    title = 'corr_' + str(i+1)
    plt.savefig(title, dpi=420, )
    plt.show()


```



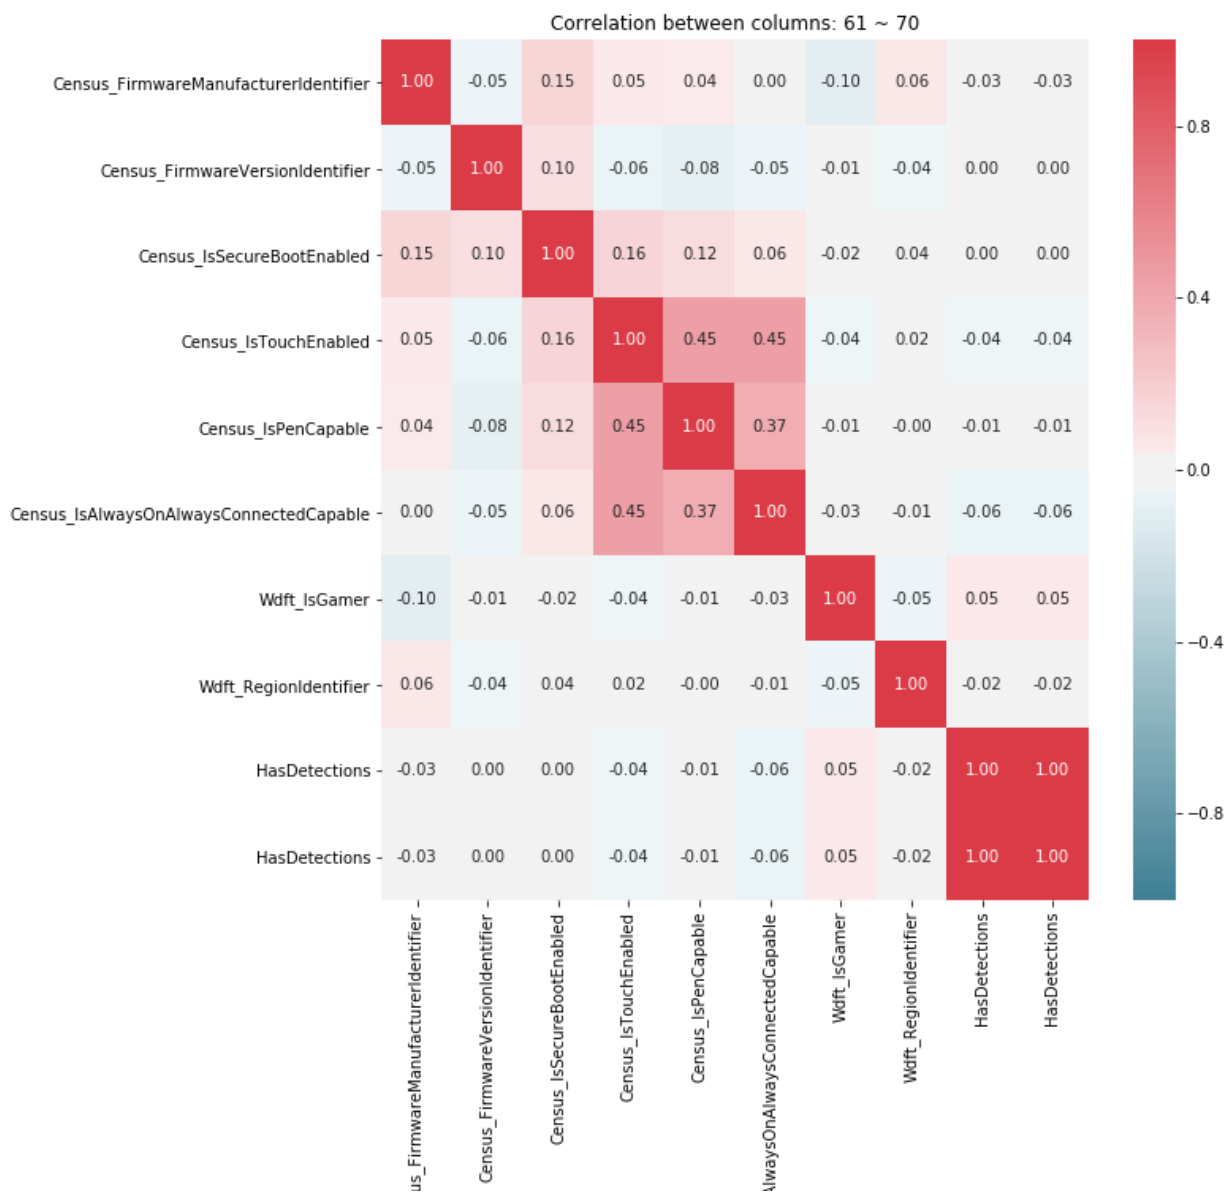












Display correlation heat maps (Correlation all together)

```
cols = train.columns
```

```
#for i in range(0, len(cols), 10): co_cols = cols
```

```
#co_cols.append(pd.Index('HasDetections')) plt.figure(figsize=(10,10))
sns.heatmap(train[co_cols].corr(), cmap='RdBu_r', annot=True, center=0.0) plt.title('Correlation
between columns: %d ~ %d' % (i, i+10)) plt.show()
```

```
In [302]: # Remove highly correlated data

corr_thresh = .99
cols = train.columns
corr = train.corr()
new_bad = []

for i in range(0, corr.shape[1]):
    for j in range(0, i):
        # Remove highly correlated features
        if (abs(corr.iloc[i, j]) >= corr_thresh):
            # Remove the feature with less information
            if (train[cols[i]].nunique() > train[cols[j]].nunique()):
                new_bad.append(cols[j])
            else:
                new_bad.append(cols[i])
```

```
In [303]: if ('HasDetections' in new_bad):
            new_bad.remove('HasDetections')

len(new_bad)
```

```
Out[303]: 2
```

```
In [304]: train.drop(new_bad, axis=1, inplace=True)
```

```
In [305]: bad_features.extend(new_bad)
bad_features = list(set(bad_features))
```

2.4 (alternative) Modified Tree Algorithm

3 Prediction

Testing trees and logistic regression with different regularizers. Then hypertuning each to select the best model.

```
In [306]: bad_features
```

```
Out[306]: ['Census_ProcessorClass',
           'AutoSampleOptIn',
           'Census_IsFlightingInternal',
           'PuaMode',
           'UacLuaenable',
           'Census_IsPortableOperatingSystem',
           'Platform',
           'Census_IsFlightsDisabled',
           'Census_IsWIMBootEnabled',
           'Census_ThresholdOptIn',
           'SMode',
           'Census_DeviceFamily',
           'Census_OSSkuName',
           'Census_IsVirtualDevice',
           'IsBeta']
```

```
In [307]: train.columns
```

```
Out[307]: Index(['ProductName', 'EngineVersion', 'AppVersion', 'AvSigVersion',
                 'RtpStateBitfield', 'IsSxsPassiveMode', 'DefaultBrowsersIdentifie
r',
                 'AVProductStatesIdentifier', 'AVProductsInstalled', 'AVProductsEna
bled',
                 'HasTpm', 'CountryIdentifier', 'CityIdentifier',
                 'OrganizationIdentifier', 'GeoNameIdentifier',
                 'LocaleEnglishNameIdentifier', 'Processor', 'OsVer', 'OsBuild',
                 'OsSuite', 'OsPlatformSubRelease', 'OsBuildLab', 'SkuEdition',
                 'IsProtected', 'IeVerIdentifier', 'SmartScreen', 'Firewall',
                 'Census_MDC2FormFactor', 'Census_OEMNameIdentifier',
                 'Census_OEMModelIdentifier', 'Census_ProcessorCoreCount',
                 'Census_ProcessorManufacturerIdentifier',
                 'Census_ProcessorModelIdentifier', 'Census_PrimaryDiskTotalCapacit
y',
                 'Census_PrimaryDiskTypeName', 'Census_SystemVolumeTotalCapacity',
                 'Census_HasOpticalDiskDrive', 'Census_TotalPhysicalRAM',
                 'Census_ChassisTypeName',
                 'Census_InternalPrimaryDiagonalDisplaySizeInInches',
                 'Census_InternalPrimaryDisplayResolutionHorizontal',
                 'Census_InternalPrimaryDisplayResolutionVertical',
                 'Census_PowerPlatformRoleName', 'Census_InternalBatteryType',
                 'Census_InternalBatteryNumberOfCharges', 'Census_OSVersion',
                 'Census_OSArchitecture', 'Census_OSBranch', 'Census_OSBuildNumbe
r',
                 'Census_OSBuildRevision', 'Census_OSEdition',
                 'Census_OSInstallTypeName', 'Census_OSInstallLanguageIdentifier',
                 'Census_OSUILocaleIdentifier', 'Census_OSWUAutoUpdateOptionsName',
                 'Census_GenuineStateName', 'Census_ActivationChannel',
                 'Census_FlightRing', 'Census_FirmwareManufacturerIdentifier',
                 'Census_FirmwareVersionIdentifier', 'Census_IsSecureBootEnabled',
                 'Census_IsTouchEnabled', 'Census_IsPenCapable',
                 'Census_IsAlwaysOnAlwaysConnectedCapable', 'Wdft_IsGamer',
                 'Wdft_RegionIdentifier', 'HasDetections'],
                dtype='object')
```

```
In [308]: len(train.columns)
```

```
Out[308]: 67
```

```
In [309]: # Setting up variables
UB = 8000
UB1 = 9000

X = train.drop('HasDetections', axis=1).values[:UB]
y = train['HasDetections'].tolist()[:UB]

X_val = train.drop('HasDetections', axis=1).values[UB:UB1]
y_val = train['HasDetections'].tolist()[UB:UB1]

X_test = train.drop('HasDetections', axis=1).values[UB1:]
y_test = train['HasDetections'].tolist()[UB1:]
```

3.1 Decision Tree Classifier

```
In [310]: from sklearn.tree import DecisionTreeClassifier
```

```
In [311]: tree_clf = DecisionTreeClassifier()
tree_clf.fit(X, y)

print("Train Acc %.3f" % (np.sum([tree_clf.predict(X) == y], axis=1) / len(X)))
print("Val Acc %.3f" % (np.sum([tree_clf.predict(X_val) == y_val], axis=1) / len(X_val)))

Train Acc 1.000
Val Acc 0.520
```

```
In [312]: for x,i in enumerate([10,10,10]):
           print(x,i)
```

```
0 10
1 10
2 10
```

3.1.2 Decision Tree: Max log2 Features

```
In [333]: # Find best decision tree with log2 max features
max_lens = list(range(1,60,1)) #[5, 10, 15, 20, 25, 30, 35, 40]
tree_log2_models = []

for i, m_len in enumerate(max_lens):
    tree_log2_clf = DecisionTreeClassifier(max_depth=m_len, max_features='1
    tree_log2_clf.fit(X, y)

    # Prediction acc
    train_acc = np.sum([tree_log2_clf.predict(X) == y], axis=1) / len(y)
    val_acc = np.sum([tree_log2_clf.predict(X_val) == y_val], axis=1) / len
    test_acc = np.sum([tree_log2_clf.predict(X_test) == y_test], axis=1) /

    # Save model and acc
    tree_log2_models.append((tree_log2_clf, train_acc[0], val_acc[0], test_
```




```
In [334]: xaxis_data = max_lens
yaxis_data_train = [i[1] for i in tree_log2_models]
yaxis_data_val = [i[2] for i in tree_log2_models]

figure(num=None, figsize=(8, 6), dpi=80, facecolor='w', edgecolor='k')

plt.plot(xaxis_data, yaxis_data_train, color='blue', label='Train Accuracy')
plt.plot(xaxis_data, yaxis_data_val, color='orange', label='Validation Accuracy')

# Plot text
plt.legend()
plt.title('Log2 Max Features Decision Tree Tuning: Max Depth')
plt.xlabel('Max Tree Depth')
plt.ylabel('Accuracy')

plt.savefig('Log2-CV.png', dpi=420)
plt.show()
```



```
In [335]: max([i[2] for i in tree_log2_models])
```

```
Out[335]: 0.609
```

```
In [336]: tree_log2_models
```

```
Out[336]: [(DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
1,
        max_features='log2', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
        splitter='best'), 0.56, 0.573, 0.5642926829268292),
 (DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
2,
        max_features='log2', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
        splitter='best'), 0.53775, 0.539, 0.5333658536585366),
 (DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
3,
        max_features='log2', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
        splitter='best'), 0.53775, 0.539, 0.5333658536585366)]
```

```
In [337]: feature_importance = list(tree_log2_models[6][0].feature_importances_)
arr = []
for i,j in zip(train.drop('HasDetections', axis=1).columns, feat_importance):
    if j > 0:
        arr.append((i,j))
arr
```

```
Out[337]: [('AppVersion', 0.005411327669183273),
 ('AvSigVersion', 0.12818671910397597),
 ('AVProductStatesIdentifier', 0.019625544389855067),
 ('AVProductsInstalled', 0.16069486380517994),
 ('HasTpm', 0.005490099761724008),
 ('CityIdentifier', 0.01065026768708279),
 ('OrganizationIdentifier', 0.002785002849290186),
 ('GeoNameIdentifier', 0.006808097717126339),
 ('OsBuildLab', 0.007049092326582092),
 ('SmartScreen', 0.5085008289759555),
 ('Census_OEMNameIdentifier', 0.021666999426615246),
 ('Census_ProcessorModelIdentifier', 0.015092964290409032),
 ('Census_PrimaryDiskTotalCapacity', 0.007871022064266582),
 ('Census_ChassisTypeName', 0.007338133670530459),
 ('Census_InternalPrimaryDiagonalDisplaySizeInInches', 0.072261690680889
3),
 ('Census_OSArchitecture', 0.004870194902264946),
 ('Census_OSBuildRevision', 0.0043459161036154895),
 ('Census_FirmwareVersionIdentifier', 0.0057440470328050195),
 ('Census_IsSecureBootEnabled', 0.0056071875426488005)]
```

```
In [338]: tree_log2_models[4]
```

```
Out[338]: (DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=5,
                                   max_features='log2', max_leaf_nodes=None,
                                   min_impurity_decrease=0.0, min_impurity_split=None,
                                   min_samples_leaf=1, min_samples_split=2,
                                   min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                   splitter='best'), 0.564375, 0.599, 0.5500731707317074)
```

```
In [318]: feat_importance = tree_log2_models[4][0].feature_importances_
```

```
In [319]: arr = []
for i,j in zip(train.drop('HasDetections', axis=1).columns, feat_importance):
    arr.append((i,j))
arr
```

```
Out[319]: [('ProductName', 0.0),
('EngineVersion', 0.0),
('AppVersion', 0.005411327669183273),
('AvSigVersion', 0.12818671910397597),
('RtpStateBitfield', 0.0),
('IsSxsPassiveMode', 0.0),
('DefaultBrowsersIdentifier', 0.0),
('AVProductStatesIdentifier', 0.019625544389855067),
('AVProductsInstalled', 0.16069486380517994),
('AVProductsEnabled', 0.0),
('HasTpm', 0.005490099761724008),
('CountryIdentifier', 0.0),
('CityIdentifier', 0.01065026768708279),
('OrganizationIdentifier', 0.002785002849290186),
('GeoNameIdentifier', 0.006808097717126339),
('LocaleEnglishNameIdentifier', 0.0),
('Processor', 0.0),
('OsVer', 0.0),
('OsBuild', 0.0),
('OsSuite', 0.0),
('OsPlatformSubRelease', 0.0),
('OsBuildLab', 0.007049092326582092),
('SkuEdition', 0.0),
('IsProtected', 0.0),
('IeVerIdentifier', 0.0),
('SmartScreen', 0.5085008289759555),
('Firewall', 0.0),
('Census_MDC2FormFactor', 0.0),
('Census_OEMNameIdentifier', 0.021666999426615246),
('Census_OEMModelIdentifier', 0.0),
('Census_ProcessorCoreCount', 0.0),
('Census_ProcessorManufacturerIdentifier', 0.0),
('Census_ProcessorModelIdentifier', 0.015092964290409032),
('Census_PrimaryDiskTotalCapacity', 0.007871022064266582),
('Census_PrimaryDiskTypeName', 0.0),
('Census_SystemVolumeTotalCapacity', 0.0),
('Census_HasOpticalDiskDrive', 0.0),
('Census_TotalPhysicalRAM', 0.0),
('Census_ChassisTypeName', 0.007338133670530459),
('Census_InternalPrimaryDiagonalDisplaySizeInInches', 0.072261690680889
3),
('Census_InternalPrimaryDisplayResolutionHorizontal', 0.0),
('Census_InternalPrimaryDisplayResolutionVertical', 0.0),
('Census_PowerPlatformRoleName', 0.0),
('Census_InternalBatteryType', 0.0),
('Census_InternalBatteryNumberOfCharges', 0.0),
('Census_OSVersion', 0.0),
('Census_OSArchitecture', 0.004870194902264946),
('Census_OSBranch', 0.0),
('Census_OSBuildNumber', 0.0),
('Census_OSBuildRevision', 0.0043459161036154895),
('Census_OSEdition', 0.0),
```

```
( 'Census_OSInstallTypeName', 0.0),
( 'Census_OSInstallLanguageIdentifier', 0.0),
( 'Census_OSUILocaleIdentifier', 0.0),
( 'Census_OSWUAutoUpdateOptionsName', 0.0),
( 'Census_GenuineStateName', 0.0),
( 'Census_ActivationChannel', 0.0),
( 'Census_FlightRing', 0.0),
( 'Census_FirmwareManufacturerIdentifier', 0.0),
( 'Census_FirmwareVersionIdentifier', 0.0057440470328050195),
( 'Census_IsSecureBootEnabled', 0.0056071875426488005),
( 'Census_IsTouchEnabled', 0.0),
( 'Census_IsPenCapable', 0.0),
( 'Census_IsAlwaysOnAlwaysConnectedCapable', 0.0),
( 'Wdft_IsGamer', 0.0),
( 'Wdft_RegionIdentifier', 0.0)]
```

```
In [339]: tree_sqrt_models = []

for i, m_len in enumerate(max_lens):
    tree_sqrt_clf = DecisionTreeClassifier(max_depth=m_len, max_features='sqrt')
    tree_sqrt_clf.fit(X, y)

    # Prediction acc
    train_acc = np.sum([tree_sqrt_clf.predict(X) == y], axis=1) / len(y)
    val_acc = np.sum([tree_sqrt_clf.predict(X_val) == y_val], axis=1) / len(y_val)
    test_acc = np.sum([tree_sqrt_clf.predict(X_test) == y_test], axis=1) / len(y_test)

    # Save model and acc
    tree_sqrt_models.append((tree_sqrt_clf, train_acc[0], val_acc[0], test_acc[0]))
```

```
In [340]: xaxis_data = max_lens
yaxis_sqrt_train = [i[1] for i in tree_sqrt_models]
yaxis_sqrt_val = [i[2] for i in tree_sqrt_models]


figure(num=None, figsize=(8, 6), dpi=80, facecolor='w', edgecolor='k')

# Plot Data
plt.plot(xaxis_data, yaxis_sqrt_train, color='blue', label="Train Accuracy")
plt.plot(xaxis_data, yaxis_sqrt_val, color='orange', label="Validation Accuracy")

# Plot text
plt.legend()
plt.title('Sqrt Max Features Decision Tree Tuning: Max Depth')
plt.xlabel('Max Tree Depth')
plt.ylabel('Accuracy')

plt.savefig('Sqrt-CV.png', dpi=420)

plt.show()
```



```
In [341]: max([i[2] for i in tree_sqrt_models])
```

```
Out[341]: 0.605
```

In [342]: tree_sqrt_models

```
Out[342]: [(DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
1,
        max_features='sqrt', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
        splitter='best'), 0.552, 0.55),
 (DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
2,
        max_features='sqrt', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
        splitter='best'), 0.54, 0.536),
 (DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
3,
        max_features='sqrt', max_leaf_nodes=None,
        min_impurity_decrease=0.0, min_impurity_split=None,
        min_samples_leaf=1, min_samples_split=2,
        min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
        splitter='best'), 0.53, 0.53)]
```

In [348]: feat_importance = tree_sqrt_models[7][0].feature_importances_

```
In [349]: arr = []
for i,j in zip(train.drop('HasDetections', axis=1).columns, feat_importance):
    if j > 0:
        arr.append((i,j))
arr
```

```
Out[349]: [('EngineVersion', 0.09797331267045815),
('AppVersion', 0.019741279003257204),
('AvSigVersion', 0.019891061703133407),
('DefaultBrowsersIdentifier', 0.006459270871591631),
('AVProductStatesIdentifier', 0.011409068068357932),
('AVProductsInstalled', 0.11096245124235501),
('CountryIdentifier', 0.005484038520492168),
('CityIdentifier', 0.03340663155847973),
('OrganizationIdentifier', 0.0019450357932891949),
('GeoNameIdentifier', 0.04800244363364314),
('LocaleEnglishNameIdentifier', 0.0026040478507890078),
('OsPlatformSubRelease', 0.014639508082357158),
('OsBuildLab', 0.016157268405932613),
('SkuEdition', 0.0029373228699584317),
('IeVerIdentifier', 0.02244268049646621),
('SmartScreen', 0.12120814002527844),
('Census_MDC2FormFactor', 0.00527028748848066),
('Census_OEMNameIdentifier', 0.02096356194315714),
('Census_OEMModelIdentifier', 0.02235835173129051),
('Census_ProcessorCoreCount', 0.005456866808907994),
('Census_ProcessorModelIdentifier', 0.04760404814757409),
('Census_PrimaryDiskTotalCapacity', 0.06452195009038093),
('Census_SystemVolumeTotalCapacity', 0.01932490188113224),
('Census_HasOpticalDiskDrive', 0.004126664992184298),
('Census_TotalPhysicalRAM', 0.015183749728726431),
('Census_ChassisTypeName', 0.0051817671596256),
('Census_InternalPrimaryDiagonalDisplaySizeInInches', 0.0044002119285425625),
('Census_InternalPrimaryDisplayResolutionHorizontal', 0.011180178809315813),
('Census_InternalPrimaryDisplayResolutionVertical', 0.01469501848960027),
('Census_PowerPlatformRoleName', 0.001103013920531091),
('Census_InternalBatteryNumberOfCharges', 0.007947131097114965),
('Census_OSVersion', 0.016009741886231256),
('Census_OSBuildNumber', 0.003268189394166195),
('Census_OSBuildRevision', 0.02749003348825741),
('Census_OSEdition', 0.011651108734424376),
('Census_OSInstallTypeName', 0.03220100213366562),
('Census_OSInstallLanguageIdentifier', 0.007971386436886583),
('Census_OSUILocaleIdentifier', 0.010367284419974506),
('Census_OSWUAutoUpdateOptionsName', 0.005845636681005605),
('Census_ActivationChannel', 0.0030759429592152414),
('Census_FlightRing', 0.0030789141592119512),
('Census_FirmwareManufacturerIdentifier', 0.011815157621338591),
('Census_FirmwareVersionIdentifier', 0.022411707614850052),
('Census_IsSecureBootEnabled', 0.009152645071616523),
('Census_IsTouchEnabled', 0.015835613484674178),
('Wdft_IsGamer', 0.01588270317606865),
('Wdft_RegionIdentifier', 0.019361667726009172)]
```



```
In [354]: arr.sort(key = lambda arr: arr[1])
```

```
In [355]: arr
```

```
Out[355]: [('Census_PowerPlatformRoleName', 0.001103013920531091),
 ('OrganizationIdentifier', 0.0019450357932891949),
 ('LocaleEnglishNameIdentifier', 0.0026040478507890078),
 ('SkuEdition', 0.0029373228699584317),
 ('Census_ActivationChannel', 0.0030759429592152414),
 ('Census_FlightRing', 0.0030789141592119512),
 ('Census_OSBuildNumber', 0.003268189394166195),
 ('Census_HasOpticalDiskDrive', 0.004126664992184298),
 ('Census_InternalPrimaryDiagonalDisplaySizeInInches', 0.0044002119285425
625),
 ('Census_ChassisTypeName', 0.0051817671596256),
 ('Census_MDC2FormFactor', 0.00527028748848066),
 ('Census_ProcessorCoreCount', 0.005456866808907994),
 ('CountryIdentifier', 0.005484038520492168),
 ('Census_OSWUAutoUpdateOptionsName', 0.005845636681005605),
 ('DefaultBrowsersIdentifier', 0.006459270871591631),
 ('Census_InternalBatteryNumberOfCharges', 0.007947131097114965),
 ('Census_OSInstallLanguageIdentifier', 0.007971386436886583),
 ('Census_IsSecureBootEnabled', 0.009152645071616523),
 ('Census_OSUILocaleIdentifier', 0.010367284419974506),
 ('Census_InternalPrimaryDisplayResolutionHorizontal', 0.0111801788093158
13),
 ('AVProductStatesIdentifier', 0.011409068068357932),
 ('Census_OSEdition', 0.011651108734424376),
 ('Census_FirmwareManufacturerIdentifier', 0.011815157621338591),
 ('OsPlatformSubRelease', 0.014639508082357158),
 ('Census_InternalPrimaryDisplayResolutionVertical', 0.0146950184896002
7),
 ('Census_TotalPhysicalRAM', 0.015183749728726431),
 ('Census_IsTouchEnabled', 0.015835613484674178),
 ('Wdft_IsGamer', 0.01588270317606865),
 ('Census_OSVersion', 0.016009741886231256),
 ('OsBuildLab', 0.016157268405932613),
 ('Census_SystemVolumeTotalCapacity', 0.01932490188113224),
 ('Wdft_RegionIdentifier', 0.019361667726009172),
 ('AppVersion', 0.019741279003257204),
 ('AvSigVersion', 0.019891061703133407),
 ('Census_OEMNameIdentifier', 0.02096356194315714),
 ('Census_OEMModelIdentifier', 0.02235835173129051),
 ('Census_FirmwareVersionIdentifier', 0.022411707614850052),
 ('IeVerIdentifier', 0.02244268049646621),
 ('Census_OSBuildRevision', 0.02749003348825741),
 ('Census_OSInstallTypeName', 0.03220100213366562),
 ('CityIdentifier', 0.03340663155847973),
 ('Census_ProcessorModelIdentifier', 0.04760404814757409),
 ('GeoNameIdentifier', 0.04800244363364314),
 ('Census_PrimaryDiskTotalCapacity', 0.06452195009038093),
 ('EngineVersion', 0.09797331267045815),
 ('AVProductsInstalled', 0.11096245124235501),
 ('SmartScreen', 0.12120814002527844)]
```

```
In [350]: sum(i[1] for i in arr)
```

```
Out[350]: 0.9999999999999999
```

```
In [ ]: [ ('AppVersion', 0.005411327669183273),
  ('AvSigVersion', 0.12818671910397597),
  ('AVProductStatesIdentifier', 0.019625544389855067),
  ('AVProductsInstalled', 0.16069486380517994),
  ('HasTpm', 0.005490099761724008),
  ('CityIdentifier', 0.01065026768708279),
  ('OrganizationIdentifier', 0.002785002849290186),
  ('GeoNameIdentifier', 0.006808097717126339),
  ('OsBuildLab', 0.007049092326582092),
  ('SmartScreen', 0.5085008289759555),
  ('Census_OEMNameIdentifier', 0.021666999426615246),
  ('Census_ProcessorModelIdentifier', 0.015092964290409032),
  ('Census_PrimaryDiskTotalCapacity', 0.007871022064266582),
  ('Census_ChassisTypeName', 0.007338133670530459),
  ('Census_InternalPrimaryDiagonalDisplaySizeInInches', 0.0722616906808893),
  ('Census_OSArchitecture', 0.004870194902264946),
  ('Census_OSBuildRevision', 0.0043459161036154895),
  ('Census_FirmwareVersionIdentifier', 0.0057440470328050195),
  ('Census_IsSecureBootEnabled', 0.0056071875426488005)]
```

```
In [217]: tree_sqrt_models[8]
```

```
Out[217]: (DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=9,
  max_features='sqrt', max_leaf_nodes=None,
  min_impurity_decrease=0.0, min_impurity_split=None,
  min_samples_leaf=1, min_samples_split=2,
  min_weight_fraction_leaf=0.0, presort=False, random_state=None,
  splitter='best'), 0.65575, 0.584)
```

```
In [218]: tree_sqrt_models[6]
```

```
Out[218]: (DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=7,
  max_features='sqrt', max_leaf_nodes=None,
  min_impurity_decrease=0.0, min_impurity_split=None,
  min_samples_leaf=1, min_samples_split=2,
  min_weight_fraction_leaf=0.0, presort=False, random_state=None,
  splitter='best'), 0.632625, 0.5795)
```

3.2 Logistic Regression

```
In [265]: from sklearn.linear_model import LogisticRegression
```

```
In [266]: ones = np.ones(X.shape[0]).reshape(-1, 1)
logit_X = np.append(X, ones, 1)

ones = np.ones(X_val.shape[0]).reshape(-1, 1)
logit_X_val = np.append(X_val, ones, 1)
```

```
In [267]: lr = LogisticRegression(penalty='l2', solver='liblinear')
lr.fit(X, y)
print("Train Acc %.4f" % (np.sum([lr.predict(X) == y], axis=1) / len(y)))
print("Test Acc %.4f" % (np.sum([lr.predict(X_val) == y_val], axis=1) / len(y_val)))
```

Train Acc 0.5270

Test Acc 0.5360

3.2.1 L2 Regularized Logistic Regression

```
In [268]: for i, c in enumerate(C):
          print(i,c)
```

```
0 0.0001
1 0.001
2 0.005
3 0.01
4 1
5 10
6 100
7 500
8 1000
9 5000
10 10000
```

```
In [277]: # Logistic regression with L2 regularization
C = [.0001, .00025, .0005, .00075, .001, .0025, .005, .0075, .01, 1, 10, 100]
#C = [.000001, .01, 1, 10, 100, 1000000]

l2_logit_models = []

for i, c in enumerate(C):
    lr = LogisticRegression(penalty='l2', C=c, solver='liblinear')
    lr.fit(X, y)

    # Prediction acc
    train_acc = np.sum([lr.predict(X) == y], axis=1) / len(y)
    val_acc = np.sum([lr.predict(X_val) == y_val], axis=1) / len(y_val)

    # Save model and acc
    l2_logit_models.append((lr, train_acc[0], val_acc[0], ))
```


```
In [278]: xaxis_data = C
yaxis_l2_logit_train = [i[1] for i in l2_logit_models]
yaxis_l2_logit_val = [i[2] for i in l2_logit_models]

figure(num=None, figsize=(8, 6), dpi=80, facecolor='w', edgecolor='k')

# Plot Data
plt.plot(np.log(xaxis_data), yaxis_l2_logit_train, color='blue', label="Train Accuracy")
plt.plot(np.log(xaxis_data), yaxis_l2_logit_val, color='orange', label="Validation Accuracy")

# Plot text
plt.legend()
plt.title('L2 Logistic Regression Tuning: Regularization Constant (Lambda)')
plt.xlabel('Log of Regularization Constant')
plt.ylabel('Accuracy')

plt.show()
```



```
In [275]: C
```

```
Out[275]: [1e-06, 0.01, 1, 10, 100, 1000000]
```

3.2.1 L1 Regularized Logistic Regression

```
In [242]: # Logistic regression with L1 regularization
C = [.0001, .001, .005, .01, 1, 10, 100, 500, 1000, 5000, 10000]

l1_logit_models = []

for i, c in enumerate(C):
    lr = LogisticRegression(penalty='l1', C=c, solver='liblinear')
    lr.fit(logit_X, y)

    # Prediction acc
    train_acc = np.sum([lr.predict(logit_X) == y], axis=1) / len(y)
    val_acc = np.sum([lr.predict(logit_X_val) == y_val], axis=1) / len(y_val)

    # Save model and acc
    l1_logit_models.append((lr, train_acc[0], val_acc[0], ))
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
31: ConvergenceWarning: Liblinear failed to converge, increase the number
of iterations.
```

```
"the number of iterations.", ConvergenceWarning)
```

```
/anaconda3/envs/py3-env/lib/python3.7/site-packages/sklearn/svm/base.py:9
```

31: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.

"the number of iterations.", ConvergenceWarning)

```
In [246]: xaxis_data = C
yaxis_l1_logit_train = [i[1] for i in l1_logit_models]
yaxis_l1_logit_val = [i[2] for i in l1_logit_models]

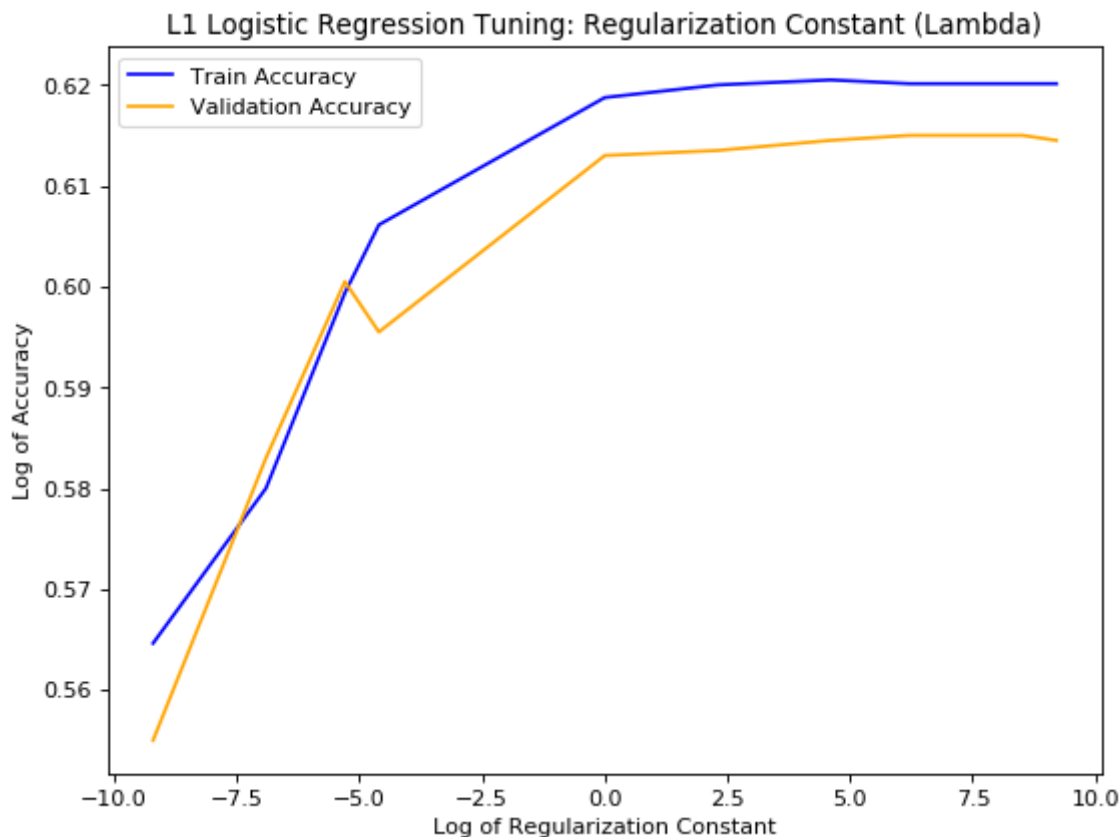
figure(num=None, figsize=(8, 6), dpi=80, facecolor='w', edgecolor='k')

# Plot Data
plt.plot(np.log(xaxis_data), yaxis_l1_logit_train, color='blue', label="Train Accuracy")
plt.plot(np.log(xaxis_data), yaxis_l1_logit_val, color='orange', label="Validation Accuracy")

# Plot text
plt.legend()
plt.title('L1 Logistic Regression Tuning: Regularization Constant (Lambda)')
plt.xlabel('Log of Regularization Constant')
plt.ylabel('Log of Accuracy')

plt.savefig('L1-Logistic-CV.png', dpi=420)

plt.show()
```



```
In [248]: max(i[2] for i in l1_logit_models)
```

```
Out[248]: 0.615
```

In [249]: l1_logit_models

```
Out[249]: [(LogisticRegression(C=0.0001, class_weight=None, dual=False,
                                fit_intercept=True, intercept_scaling=1, max_iter=100,
                                multi_class='warn', n_jobs=None, penalty='l1', random_state=None,
                                solver='liblinear', tol=0.0001, verbose=0, warm_start=False),
            0.564625,
            0.555),
            (LogisticRegression(C=0.001, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.58, 0.583),
            (LogisticRegression(C=0.005, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.599375, 0.6005),
            (LogisticRegression(C=0.01, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.606125, 0.5955),
            (LogisticRegression(C=1, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.61875, 0.613),
            (LogisticRegression(C=10, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.62, 0.6135),
            (LogisticRegression(C=100, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.6205, 0.6145),
            (LogisticRegression(C=500, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.620125, 0.615),
            (LogisticRegression(C=1000, class_weight=None, dual=False, fit_intercept=True,
                                intercept_scaling=1, max_iter=100, multi_class='warn',
                                n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                                tol=0.0001, verbose=0, warm_start=False), 0.620125, 0.615),
```

```
(LogisticRegression(C=5000, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                    tol=0.0001, verbose=0, warm_start=False), 0.620125, 0.615),
(LogisticRegression(C=10000, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, max_iter=100, multi_class='warn',
                    n_jobs=None, penalty='l1', random_state=None, solver='liblinear',
                    tol=0.0001, verbose=0, warm_start=False), 0.620125, 0.6145)]
```

```
In [245]: lr.coef_
1.34071174e-01,  9.24543856e-01, -1.17283276e-04,
1.25270217e-05, -3.73299198e-01, -4.12536719e-02,
-1.12407945e-01, -1.01844009e-05, -1.07386681e-06,
3.25142472e-03,  2.50192511e-04,  6.77801714e-04,
-4.27973916e-01, -2.59649967e-01, -1.67193104e-04,
6.36840642e-04,  1.90028855e-02,  3.93704222e-03,
-6.52900240e-02,  7.55056988e-01, -1.13340979e-03,
-3.55804477e-01,  6.40164038e-02, -1.72748133e-02,
-6.60740940e-06,  9.02197758e-07,  2.15952515e-02,
-6.86974590e-04, -3.60060366e-05,  4.30624027e-07,
3.91146735e-03, -3.19348600e-07,  7.75166812e-02,
1.28025299e-05, -1.72649144e-02,  6.85152009e-04,
7.65999731e-05, -2.32659264e-04,  3.95148641e-02,
-5.89359759e-02, -3.90924937e-11, -6.83265308e-03,
1.77440627e-02,  4.73746292e-02,  2.39885697e-04,
-2.13440920e-05,  6.25937763e-02, -1.44939411e-03,
3.01828511e-03, -3.58398447e-04, -1.10869333e-02,
2.72685975e-01, -2.81175681e-02, -4.54344232e-02,
1.22789987e-04, -1.14171674e-07,  4.13567776e-02,
-3.75129640e-01,  3.35626291e-01, -3.93827940e-01,
```

```
In [ ]: print ("C = 100 is the best")
```

```
In [ ]:
```

```
In [ ]: c = 100000
lr = LogisticRegression(penalty='l1', C=c, solver='liblinear')
lr.fit(logit_X, y)

# Prediction acc
train_acc = np.sum([lr.predict(logit_X) == y], axis=1) / len(y)
val_acc = np.sum([lr.predict(logit_X_val) == y_val], axis=1) / len(y_val)

# Save model and acc
l1_logit_models.append((lr, train_acc[0], val_acc[0], ))
```

```
In [ ]:
```